

Data Wrangling

Data wrangling generally refers to the process of getting a data set ready for analysis. Why would we need to do that?

Real-world data can be messy. Data sets are recorded and assembled by humans, and humans make mistakes. A single data set might be created and updated by multiple people who may decide to do things in slightly different ways. On a spreadsheet, one person might decide to leave cells with missing data blank, another might enter "NaN", while a third may enter "missing". If the data has many rows, one person might decide to repeat the column headers partway down so they don't have to scroll up to see them. Any of these things mean that the data set cannot be analyzed "as is" and wrangling will be required.

Even in a tightly controlled laboratory setting in which data are collected via computer and automatically written out to data files, some data wrangling might be required. There might be a separate data file for each subject or experimental session, meaning that these separate files will have to be combined into a single data set before analysis.

Our main wrangling tool is pandas, so we can go ahead and import it.

```
In [1]: import pandas as pd
```

Loading

For our wrangling practice today, we'll look at a data set containing various measurements on breast cancer patients. The file is called `breast_cancer_data.csv`, and you should place it in the "data" folder you should already have in the same directory as this notebook.

Let's import it as a pandas dataframe.

```
In [2]: bcd = pd.read_csv('./data/breast_cancer_data.csv')  
bcd
```

Out [2]:

	patient_id	clump_thickness	cell_size_uniformity	cell_shape_uniformity	margin
0	1000025	5.0	1.0	1	
1	1002945	5.0	4.0	4	
2	1015425	3.0	1.0	1	
3	1016277	6.0	8.0	8	
4	1017023	4.0	1.0	1	
...	
694	776715	3.0	1.0	1	
695	841769	2.0	1.0	1	
696	888820	5.0	10.0	10	
697	897471	4.0	8.0	6	
698	897471	4.0	8.0	8	

699 rows × 12 columns

Before we do any actual wrangling, let's get familiar with the data frame in its current form.

Exploring the Data Frame

We can explore the data frame by looking at it's attributes, such as its shape, column names, and data types:

In [3]: `bcd.columns`

Out [3]: `Index(['patient_id', 'clump_thickness', 'cell_size_uniformity', 'cell_shape_uniformity', 'marginal_adhesion', 'single_ep_cell_size', 'bare_nuclei', 'bland_chromatin', 'normal_nucleoli', 'mitoses', 'class', 'doctor_name'], dtype='object')`

Use the cells below to get the shape and data types (`dtypes`) of our data frame.

In [4]: `bcd.shape`

Out [4]: `(699, 12)`

In [5]: `bcd.dtypes`

```
Out[5]: patient_id          int64
         clump_thickness    float64
         cell_size_uniformity float64
         cell_shape_uniformity int64
         marginal_adhesion   int64
         single_ep_cell_size int64
         bare_nuclei          object
         bland_chromatin     float64
         normal_nucleoli     float64
         mitoses              int64
         class                object
         doctor_name          object
         dtype: object
```

In the cell below, use the `describe()` method to get a summary of the numerical columns.

```
In [6]: bcd.describe()
```

```
Out[6]:      patient_id  clump_thickness  cell_size_uniformity  cell_shape_uniformity  m
count  6.990000e+02        698.000000        698.000000        699.000000
mean   1.071704e+06       4.416905       3.137536       3.207439
std    6.170957e+05       2.817673       3.052575       2.971913
min    6.163400e+04       1.000000       1.000000       1.000000
25%    8.706885e+05       2.000000       1.000000       1.000000
50%    1.171710e+06       4.000000       1.000000       1.000000
75%    1.238298e+06       6.000000       5.000000       5.000000
max    1.345435e+07      10.000000      10.000000      10.000000
```

Modifying a text column

We'll often want to "tune up" columns that contain text. We might encounter, for example, a column containing full names that we need to break up into separate columns for the first and last names.

Let's look at the column for the doctors' names. Use the cell below to take a peek.

```
In [8]: bcd['doctor_name']
```

```
Out[8]: 0      Dr. Doe
1      Dr. Smith
2      Dr. Lee
3      Dr. Smith
4      Dr. Wong
...
694     Dr. Lee
695     Dr. Smith
696     Dr. Lee
697     Dr. Lee
698     Dr. Wong
Name: doctor_name, Length: 699, dtype: object
```

The doctors' name data are redundant; each one has a "Dr. " in front of the actual name, but we already know these are doctors by the column name. Further, the entries have white space in them, which can cause us problems down the road. So let's modify this column so it only contains the surnames of the doctors.

One great thing about pandas is that it has versions of many of Python's string methods that operate *element-wise on an entire column of strings*. Here, we want to separate the "Dr. " from the actual name, which is exactly what Python's `str.split()` function does. So chances are, pandas has a version of this function that operates element-wise on data frames.

String Splitting Review:

Let's briefly remind ourselves of splitting up Python strings and extracting bits of them.

```
In [9]: # Here's a string of the form: surname, first initial.
myStr = 'SirString, A.'
print(myStr)
```

SirString, A.

Let's say we wanted to get the surname. We could split this string into a Python list at the white space like this:

```
In [10]: spltStr = myStr.split()    # split() defaults to splitting at white space
print(spltStr)
```

['SirString', 'A.]

We now have a list in which the items contain the text on either side of the split. This is close to what we want: the first entry in the list has the surname, but it also has an unwanted comma.

Let's split the string at the comma instead:

```
In [11]: spltStr = myStr.split(',') # tell Python to split at commas
print(spltStr)

['SirString', ' A.]
```

Now we have isolated the last name, and we can fetch it by indexing:

```
In [12]: surname = spltStr[0]
print(surname)

SirString
```

In the cell below, see if you can extract the surname from `myStr` in one line of code:

```
In [13]: myStr.split(',')[0]

Out[13]: 'SirString'
```

Alright, time to replace the `bcd['doctor_name']` column values with just the doctors' last names.

We could do this in one step, but let's break it out for clarity. First, let's copy the name column out into a new series.

```
In [14]: dr_names = bcd['doctor_name']
dr_names

Out[14]: 0      Dr. Doe
1      Dr. Smith
2      Dr. Lee
3      Dr. Smith
4      Dr. Wong
...
694     Dr. Lee
695     Dr. Smith
696     Dr. Lee
697     Dr. Lee
698     Dr. Wong
Name: doctor_name, Length: 699, dtype: object
```

***Note:** pandas objects behave like ordinary Python objects. So, strictly speaking, we have not created a new object (pandas Series), rather, we have created a new label that refers to the "doctor_name" column of `bcd`.

In the cell below, use the `id()` function to compare the object IDs of `dr_names` and the corresponding column of `bcd`.

```
In [20]: print(id(bcd['doctor_name']), id(dr_names))
```

```
5513206736 5513206736
```

Now let's split all the names in the `doctor_name` column at the whitespace by using pandas `DataFrame.str.split()` function.

```
In [16]: split_dr_names = dr_names.str.split()  
split_dr_names
```

```
Out[16]: 0      [Dr., Doe]  
1      [Dr., Smith]  
2      [Dr., Lee]  
3      [Dr., Smith]  
4      [Dr., Wong]  
     ...  
694     [Dr., Lee]  
695     [Dr., Smith]  
696     [Dr., Lee]  
697     [Dr., Lee]  
698     [Dr., Wong]  
Name: doctor_name, Length: 699, dtype: object
```

`DataFrame.str.split()`, however, does create a new object.

Use the cell below to confirm that the `split()` spawned a new object.

```
In [21]: print(id(bcd['doctor_name']), id(split_dr_names))
```

```
5513206736 5212828432
```

Now we have a column of lists, each with two elements. The first element of each list is the "Dr. " bit, and the second consists of the surnames we want.

We can get these by using pandas string indexing, `Series.str[index]`.

```
In [22]: surnames = split_dr_names.str[1]  
surnames
```

```
Out[22]: 0      Doe
          1      Smith
          2      Lee
          3      Smith
          4      Wong
          ...
          694     Lee
          695     Smith
          696     Lee
          697     Lee
          698     Wong
Name: doctor_name, Length: 699, dtype: object
```

Note that, like the splitting, the string indexing worked on the entire `Series` automatically.

Now we can change the column in our main data frame, `bcd`.

```
In [23]: bcd['doctor_name'] = surnames
```

```
In [24]: bcd['doctor_name']
```

```
Out[24]: 0      Doe
          1      Smith
          2      Lee
          3      Smith
          4      Wong
          ...
          694     Lee
          695     Smith
          696     Lee
          697     Lee
          698     Wong
Name: doctor_name, Length: 699, dtype: object
```

Success!

Converting a column type (and other aggravations)

Let's look at those data types again.

```
In [25]: bcd.dtypes
```

```
Out[25]: patient_id          int64
         clump_thickness     float64
         cell_size_uniformity float64
         cell_shape_uniformity int64
         marginal_adhesion    int64
         single_ep_cell_size   int64
         bare_nuclei           object
         bland_chromatin       float64
         normal_nucleoli        float64
         mitoses                int64
         class                  object
         doctor_name            object
         dtype: object
```

Notice that "class" and "doctor_name" are of dtype "object", which refers to a general purpose column type, and is how pandas imports text columns by default. Most of the others are numeric (integers or floats), except for "bare_nuclei".

In the cell below, take a quick glance at 'bcd' again, and see if the "bare_nuclei" column should be a different data type than, say, the "marginal_adhesion" column.

```
In [26]: bcd['bare_nuclei']
```

```
Out[26]: 0      1
         1      10
         2      2
         3      4
         4      1
         ..
         694     2
         695     1
         696     3
         697     4
         698     5
Name: bare_nuclei, Length: 699, dtype: object
```

It looks like "bare_nuclei" was intended to be a numeric column, so let's try and convert it using the `DataFrame.astype()` converter method.

```
In [27]: bcd['bare_nuclei'] = bcd['bare_nuclei'].astype('int64')
```

```
-----  
ValueError                                     Traceback (most recent call last)  
Cell In[27], line 1  
----> 1 bcd['bare_nuclei'] = bcd['bare_nuclei'].astype('int64')  
  
File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/generic.py:653  
4, in NDFrame.astype(self, dtype, copy, errors)  
    6530     results = [ser.astype(dtype, copy=copy) for _, ser in self.items()  
()]  
    6532 else:  
    6533     # else, only a single dtype is given  
-> 6534     new_data = self._mgr.astype(dtype=dtype, copy=copy, errors=errors)  
    6535     res = self._constructor_from_mgr(new_data, axes=new_data.axes)  
    6536     return res.__finalize__(self, method="astype")  
  
File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/internals/managers.py:414, in BaseBlockManager.astype(self, dtype, copy, errors)  
    411 elif using_copy_on_write():  
    412     copy = False  
-> 414 return self.apply(  
    415     "astype",  
    416     dtype=dtype,  
    417     copy=copy,  
    418     errors=errors,  
    419     using_cow=using_copy_on_write(),  
    420 )  
  
File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/internals/managers.py:354, in BaseBlockManager.apply(self, f, align_keys, **kwargs)  
    352     applied = b.apply(f, **kwargs)  
    353 else:  
-> 354     applied = getattr(b, f)(**kwargs)  
    355     result_blocks = extend_blocks(applied, result_blocks)  
    357 out = type(self).from_blocks(result_blocks, self.axes)  
  
File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/internals/blocks.py:616, in Block.astype(self, dtype, copy, errors, using_cow)  
    596 """  
    597 Coerce to the new dtype.  
    598 (...)  
    612 Block  
    613 """  
    614 values = self.values  
-> 616 new_values = astype_array_safe(values, dtype, copy=copy, errors=errors)  
    618 new_values = maybe_coerce_values(new_values)  
    620 refs = None  
  
File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/dtypes/astype.py:238, in astype_array_safe(values, dtype, copy, errors)  
    235     dtype = dtype.numpy_dtype  
    237 try:  
-> 238     new_values = astype_array(values, dtype, copy=copy)  
    239 except (ValueError, TypeError):
```

```
240      # e.g. _astype_nansafe can fail on object-dtype of strings
241      # trying to convert to float
242      if errors == "ignore":
243
244          File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/dtypes/astype.py:183, in astype_array(values, dtype, copy)
245              values = values.astype(dtype, copy=copy)
246          else:
247              values = _astype_nansafe(values, dtype, copy=copy)
248          # in pandas we don't store numpy str dtypes, so convert to object
249          if isinstance(dtype, np.dtype) and issubclass(values.dtype.type, str):
250
251              File /opt/anaconda3/lib/python3.11/site-packages/pandas/core/dtypes/astype.py:134, in _astype_nansafe(arr, dtype, copy, skipna)
252                  raise ValueError(msg)
253          if copy or arr.dtype == object or dtype == object:
254              # Explicit copy, or required since NumPy can't view from / to object.
255          return arr.astype(dtype, copy=True)
256          return arr.astype(dtype, copy=copy)
257
258      ValueError: invalid literal for int() with base 10: '?'
```

And, argh, we get an error! If we look at the bottom of the error message, it seems that the error involves question marks ("?") in the data, which would also explain why this column imported as text rather than numbers in the first place.

Let's check.

In the cell below, use logical indexing to show the rows of `bcd` in which `bcd["bare_nuclei"]` contains a question mark.

```
In [28]: bcd[bcd["bare_nuclei"] == "?"]
```

Out [28]:

	patient_id	clump_thickness	cell_size_uniformity	cell_shape_uniformity	margin
23	1057013	8.0	4.0	5	
40	1096800	6.0	6.0	6	
139	1183246	1.0	1.0	1	
145	1184840	1.0	1.0	3	
158	1193683	1.0	1.0	2	
164	1197510	5.0	1.0	1	
235	1241232	3.0	1.0	4	
249	169356	3.0	1.0	1	
275	432809	3.0	1.0	3	
292	563649	8.0	8.0	8	
294	606140	1.0	1.0	1	
297	61634	5.0	4.0	3	
315	704168	4.0	6.0	5	
321	733639	3.0	1.0	1	
411	1238464	1.0	1.0	1	
617	1057067	1.0	1.0	1	

Sure enough. Rather than leaving the cells of missing values empty, somebody has made the poor decision to enter question marks instead.

When you are dealing with other peoples' data, you'll find that this sort of the happens a LOT. It can be very aggravating, so we need to learn to treat these things as challenging puzzles instead of hassles!

Let's replace the question marks with nothing, so that this column becomes consistent with the rest. Fortunately, `DataFrame` (and `Series`) objects have a `replace()` function built in, so let's use that.

In [29]: `bcd['bare_nuclei'] = bcd['bare_nuclei'].replace('?', '')`

In the cell below, confirm that we no longer have question marks in our "bare_nuclei" column.

In [30]: `bcd[bcd["bare_nuclei"] == "?"]`

```
Out[30]: patient_id  clump_thickness  cell_size_uniformity  cell_shape_uniformity  marginal_a...
```

Note: As mentioned above, extracting columns or other subsets of data from a pandas `DataFrame` or `Series` does not create a new object but rather a new label to the existing object.

So, for example, `the_IDs = bcd['patient_id']` does not make a new object, but rather creates a second label referring to the original object (consistent with the behavior of base Python).

In general, however, pandas methods (functions) *do* create new objects. Thus, the step of assigning the output of `.replace()` back to the original data frame column is necessary.

In the cells below, confirm that the output of `.replace()` and `bcd['bare_nuclei']` have different IDs.

```
In [33]: # first check for a value with which to test  
replaced = bcd['bare_nuclei'].replace(10,11)  
print(id(bcd['bare_nuclei']), id(replaced))
```

```
6068658320 6068630288
```

And now we can convert the column to numeric values.

```
In [34]: bcd['bare_nuclei'] = pd.to_numeric(bcd['bare_nuclei'])
```

In the cell below, check the data types of columns in `bcd`.

```
In [35]: bcd.dtypes
```

```
Out[35]: patient_id           int64
          clump_thickness      float64
          cell_size_uniformity float64
          cell_shape_uniformity int64
          marginal_adhesion    int64
          single_ep_cell_size   int64
          bare_nuclei           float64
          bland_chromatin       float64
          normal_nucleoli       float64
          mitoses                int64
          class                  object
          doctor_name            object
          dtype: object
```

Okay! We have now have gotten our data somewhat into shape, meaning:

- missing data are actually missing
- columns of numeric data are numeric in type
- the column of doctor names contains only last names

So now we can explore some ways to deal with missing values.

Dealing with missing data

Finding missing values

Even though this dataset isn't all that large:

```
In [36]: bcd.shape
```

```
Out[36]: (699, 12)
```

699 rows is lot to look through "by hand" in order to find missing values.

We can test for missing values using the `DataFrame.isna()` method.

```
In [37]: bcd.isna()
```

Out[37]:

	patient_id	clump_thickness	cell_size_uniformity	cell_shape_uniformity	margin:
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...
694	False	False	False	False	False
695	False	False	False	False	False
696	False	False	False	False	False
697	False	False	False	False	False
698	False	False	False	False	False

699 rows × 12 columns

By itself, that doesn't help us much. But if we combine it with summation (remember that `True` values count as 1 and `False` counts as zero):

In [38]:

```
bcd.isna().sum()
```

Out[38]:

patient_id	0
clump_thickness	1
cell_size_uniformity	1
cell_shape_uniformity	0
marginal_adhesion	0
single_ep_cell_size	0
bare_nuclei	18
bland_chromatin	4
normal_nucleoli	1
mitoses	0
class	0
doctor_name	0
dtype:	int64

Now we have the counts by variable, and can easily see that there are missing values for a few of the variables.

Let's check some of the rows with missing values and make sure everything else looks normal in those rows. Notice above that the output of `.isna()` is Boolean, so we can use it to do logical indexing.

In [39]:

```
bcd[bcd['bland_chromatin'].isna()]
```

Out[39]:

	patient_id	clump_thickness	cell_size_uniformity	cell_shape_uniformity	margin
342	814265	2.0	1.0		1
343	814911	1.0	1.0		1
359	873549	10.0	3.0		5
365	897172	2.0	1.0		1

Okay, it looks like all of the other columns look fine.

In the cell below, check the rows that have missing values for either clump thickness or cell size uniformity. Do this in one go rather than separately (remember about the element-wise or operator, "|".

In [40]:

```
bcd[(bcd["clump_thickness"].isna()) | (bcd["cell_size_uniformity"].isna())]
```

Out[40]:

	patient_id	clump_thickness	cell_size_uniformity	cell_shape_uniformity	marginal
6	1018099	1.0	NaN		1
12	1041801	NaN		3.0	3

So far so good. It looks like the rows that have missing values just have one missing value, and everything else seems fine. But let's do check that no rows have more than one missing value.

To do this, we can sum the number of missing values across the columns (i.e. within each row), and then see what the maximum number of missing values within a row is.

In [41]:

```
row_na_totals = bcd.isna().sum(axis = 1)
row_na_totals.max()
```

Out[41]: 1

So we see that no row has more than one missing value.

In the cell below, do the above calculation in one line.

In [42]:

```
print(bcd.isna().sum(axis = 1).max())
```

1

Dealing with missing values

Now that we have determined that there are missing values, we have to determine how to deal with them.

Ignoring missing values elementwise

One way to handle missing values is just to ignore them. Most of the standard math and statistical functions will do that by default.

So this:

```
In [43]: bcd['clump_thickness'].mean()
```

```
Out[43]: 4.416905444126074
```

Computes the mean clump thickness ignoring the one missing value.

We can compute the mean (again ignoring missing values) for all the numeric columns like this:

```
In [44]: bcd.mean(numeric_only = True) # the numeric_only refers to columns, not miss
```

```
Out[44]: patient_id      1.071704e+06
clump_thickness      4.416905e+00
cell_size_uniformity 3.137536e+00
cell_shape_uniformity 3.207439e+00
marginal_adhesion    2.793991e+00
single_ep_cell_size   3.216023e+00
bare_nuclei          3.538913e+00
bland_chromatin      3.447482e+00
normal_nucleoli      2.868195e+00
mitoses              1.589413e+00
dtype: float64
```

That worked, but the output is a little awkward because the patient ID is being treated as a numeric variable. We can fix that by converting the patient ID variable to a string variable.

```
In [45]: bcd['patient_id'] = bcd['patient_id'].astype('string')
```

And now the means should look a little better because we won't have the mean for the ID column in the millions>

Recompute the mean for the numeric columns in the cell below.

```
In [46]: bcd.mean(numeric_only = True)
```

```
Out[46]: clump_thickness      4.416905
          cell_size_uniformity   3.137536
          cell_shape_uniformity   3.207439
          marginal_adhesion       2.793991
          single_ep_cell_size     3.216023
          bare_nuclei              3.538913
          bland_chromatin         3.447482
          normal_nucleoli          2.868195
          mitoses                  1.589413
          dtype: float64
```

Removing missing values

We are about to start learning how to remove missing values from our data frame, however...

Before we start messing around too much with the values in our data frame, let's make sure we can easily "hit the reset button" and get back to a nice starting point. To do this, we'll want to

- reload the data
- modify the column of Dr. names
- set the patient ID to type str
- remove the question marks from the bare nuclei column
- set the bare nuclei column to numeric

This is a perfect job for a function!

In the cell below, finish writing the function to reset our data frame to the desired starting point.

```
In [71]: def hit_reset():
    bcd = pd.read_csv('./data/breast_cancer_data.csv')
    bcd['patient_id'] = bcd['patient_id'].astype(str)
    bcd['doctor_name'] = bcd['doctor_name'].str.split().str[1]
    bcd['bare_nuclei'] = bcd['bare_nuclei'].replace('?', pd.NA)
    bcd['bare_nuclei'] = pd.to_numeric(bcd['bare_nuclei'])
    return bcd
```

Removing rows with missing values

Obviously, rows in which all values are missing won't do us any good, so we can drop them with:

```
In [72]: bcd = bcd.dropna(how = 'all')
```

This drops rows in which *all* of the values are missing. This code ran without error, but we know it also didn't do anything in this case because we don't have any rows in which all the values are missing!

Sometimes a case can be made for throwing out all observations (rows) that are incomplete, that is, if they contain *any* missing values.

```
In [73]: bcd = bcd.dropna(how = 'any')
```

In the cell below, check the (new) shape of `bcd`.

```
In [74]: bcd.shape
```

```
Out[74]: (699, 7)
```

It should have fewer rows now.

And now is a perfect time to test our function! In the cell below, hit the reset button on `bcd`.

```
In [75]: bcd = hit_reset()
```

Check the shape.

```
In [76]: bcd.shape
```

```
Out[76]: (699, 12)
```

Check the data types of the columns.

```
In [77]: bcd.dtypes
```

```
Out[77]: patient_id          object
         clump_thickness      float64
         cell_size_uniformity float64
         cell_shape_uniformity int64
         marginal_adhesion    int64
         single_ep_cell_size   int64
         bare_nuclei           float64
         bland_chromatin       float64
         normal_nucleoli       float64
         mitoses                int64
         class                  object
         doctor_name            object
         dtype: object
```

Check the doctor name column.

```
In [78]: bcd['doctor_name']
```

```
Out[78]: 0      Doe
1      Smith
2      Lee
3      Smith
4      Wong
...
694    Lee
695    Smith
696    Lee
697    Lee
698    Wong
Name: doctor_name, Length: 699, dtype: object
```

Removing columns with missing values

And we could do the same for columns if we wished, though this is less frequently done.
We just need to change the axis (direction) over which `DataFrame.dropna()` works.

```
In [79]: bcd = bcd.dropna(axis = 1, how = 'any') # drop columns rather than rows
```

This leaves us with only the complete columns.

```
In [80]: bcd.shape
```

```
Out[80]: (699, 7)
```

Let's see which they are.

```
In [81]: bcd.columns
```

```
Out[81]: Index(['patient_id', 'cell_shape_uniformity', 'marginal_adhesion',
               'single_ep_cell_size', 'mitoses', 'class', 'doctor_name'],
               dtype='object')
```

Filling in missing values

Occasionally, we may want to fill in missing values. This isn't very common, but might be useful if some other function you are using doesn't handle missing values gracefully.

Before filling in missing values, we need to restore our data frame so it actually has missing values. Good thing we wrote that function!

```
In [82]: bcd = hit_reset()
```

We can fill in missing values with any single value we want, such as a zero.

```
In [83]: bcd = bcd.fillna(0)
```

In the cell below, check to see that we no longer have missing values.

```
In [84]: print(bcd.isna().sum(axis = 1).max())
```

0

In the cell below, reset the data and verify that the missing data are back.

```
In [85]: bcd = hit_reset()
print(bcd.isna().sum(axis = 1).max())
```

1

In the cell below, fill the missing values in each column with the column mean. (Hint: this is pandas, so this is actually easy!)

```
In [87]: bcd = bcd.fillna(bcd.mean(numeric_only=True))
```

And now verify that there are no more missing values.

```
In [88]: print(bcd.select_dtypes(include='number').isna().sum(axis=1).max())
```

0

Summary

In this tutorial, we learned or remembered how to do some of the foundational data wrangling tasks. These are:

- importing data into pandas from a data file

- cleaning up the data in the columns
- converting columns to the appropriate type
- removing or filling in missing values